

# Similar images with Product Quantization

Mathias Chouilly, Lefty

February 24, 2016

## Introduction

Goal

Constraints

How to start ?

## Product Quantization

Original Version (Jégou 2011)

Multi-Index (Yandex 2012)

Optimized Product Quantization (Microsoft 2013)

Fast Bilayer Product Quantization (Yandex 2014)

## Implementation details

Lefty Product Quantization

Libraries

Machines and Servers

## Conclusion



# Goal

- Given an image, find similar image in our dataset

# Goal

- Given an image, find similar image in our dataset
- Mathy version: given a vector, find its nearest neighbors

# Constraints

- Handle very big dataset (possibly 1B elements)
- Fast (Ideally  $< 1s$ )
- Precise (Image should be similar to query)

## How to start ?

- Image to Vector: we already use GoogLeNet features (1024 floats) for image classification

## How to start ?

- Image to Vector: we already use GoogLeNet features (1024 floats) for image classification
- GoogLeNet is a Deep Neural Network that achieved state-of-the-art on ImageNet 2014 competition

## How to start ?

- Image to Vector: we already use GoogLeNet features (1024 floats) for image classification
- GoogLeNet is a Deep Neural Network that achieved state-of-the-art on ImageNet 2014 competition
- It is very fast to compute compared to similar performing models



## How to start ?

- Image to Vector: we already use GoogLeNet features (1024 floats) for image classification
- GoogLeNet is a Deep Neural Network that achieved state-of-the-art on ImageNet 2014 competition
- It is very fast to compute compared to similar performing models
- Let's use that

## Original Version (Jégou 2011)

- Decompose the space into a Cartesian product of low dimensional subspaces
- Quantize each subspace separately
- A vector is represented by a short code composed of its subspace quantization indices.

## Original Version (Jégou 2011)

- Decompose the space into a Cartesian product of low dimensional subspaces
- Quantize each subspace separately
- A vector is represented by a short code composed of its subspace quantization indices.
- The euclidean distance between two vectors can be efficiently estimated from their codes.
- An asymmetric distance calculation scheme increases the precision by computing the approximate distance between a vector and a code.

## Practical use

- Use a coarse quantizer (k-means)  $q_c$  to build an inverted index in original space

## Practical use

- Use a coarse quantizer (k-means)  $q_c$  to build an inverted index in original space
- Each vector  $y$  is associated to its closest centroid  $q_c(y)$ , and we denote the residual vector:  $r(y) = y - q_c(y)$

## Practical use

- Use a coarse quantizer (k-means)  $q_c$  to build an inverted index in original space
- Each vector  $y$  is associated to its closest centroid  $q_c(y)$ , and we denote the residual vector:  $r(y) = y - q_c(y)$
- Residual vectors are product quantized by  $q_p$ , such that each vector is approximated by  $y \approx q_c(y) + q_p(r(y))$

## Practical use

- Use a coarse quantizer (k-means)  $q_c$  to build an inverted index in original space
- Each vector  $y$  is associated to its closest centroid  $q_c(y)$ , and we denote the residual vector:  $r(y) = y - q_c(y)$
- Residual vectors are product quantized by  $q_p$ , such that each vector is approximated by  $y \approx q_c(y) + q_p(r(y))$
- Given a query vector  $x$ , find its nearest coarse centroids, compute residuals to each of these coarse centroids

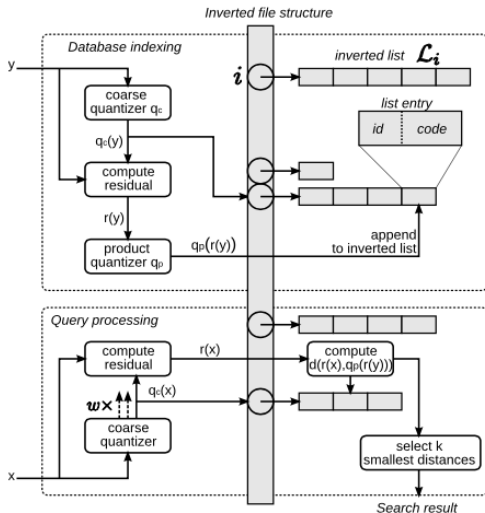
## Practical use

- Use a coarse quantizer (k-means)  $q_c$  to build an inverted index in original space
- Each vector  $y$  is associated to its closest centroid  $q_c(y)$ , and we denote the residual vector:  $r(y) = y - q_c(y)$
- Residual vectors are product quantized by  $q_p$ , such that each vector is approximated by  $y \approx q_c(y) + q_p(r(y))$
- Given a query vector  $x$ , find its nearest coarse centroids, compute residuals to each of these coarse centroids
- Use the quantization codes to compute distances and find the smallest






## Practical use



## Multi-Index (Yandex 2012)



- Build on Jégou work
- Replace the coarse quantizer by a product quantizer
- Use multi-sequence algorithm to sort coarse centroids
- Space can be divided into  $K^2$  cells, just using  $2K$  centroids
- The finer division of the space implies a higher retrieval accuracy, for the same list length as in the simple index case

# Optimized Product Quantization (Microsoft 2013)

- Replace Product Quantization by Optimized Product Quantization
- Find an optimal space decomposition when choosing low dimensional subspaces

# Optimized Product Quantization (Microsoft 2013)

- Replace Product Quantization by Optimized Product Quantization
- Find an optimal space decomposition when choosing low dimensional subspaces
- Compute covariance matrix of the data
- Use eigen value allocation algorithm to build a rotation matrix
- Do Product Quantization on the rotated data

# Fast Bilayer Product Quantization (Yandex 2014)

- Use the assumption that all residuals are compressed in the same space
- Precompute the norms at launch
- Precompute scalar products at each query
- Use precomputed tables to compute distances very fast



## Lefty Product Quantization

- Use Optimized Product Quantization as an indexing-level quantizer with 16k centroids in each subspace of 512 dimensions. That is 256M possible centroids.
- Use Optimized Product Quantization as a compressing-level quantizer with 256 centroids in each subspace of 32 dimensions. That is  $10^{77}$  possible centroids.
- Thus each vector is associated with its indexing ids (2 int) and its compressing ids (32 bytes)

## Lefty Product Quantization

- Use Optimized Product Quantization as an indexing-level quantizer with 16k centroids in each subspace of 512 dimensions. That is 256M possible centroids.
- Use Optimized Product Quantization as a compressing-level quantizer with 256 centroids in each subspace of 32 dimensions. That is  $10^{77}$  possible centroids.
- Thus each vector is associated with its indexing ids (2 int) and its compressing ids (32 bytes)
- Precompute the norms and scalar products to speed up distances calculations (from 1024 multiplications and 1023 additions to 96 additions, for each candidate)



## Indexing process

- Rotate features, and quantize them with indexing-level quantizer -> 2 indexing ids

## Indexing process

- Rotate features, and quantize them with indexing-level quantizer -> 2 indexing ids
- Compute residuals, rotate them, and quantize them with compressing-level quantizer -> 32 compressing ids

## Indexing process

- Rotate features, and quantize them with indexing-level quantizer -> 2 indexing ids
- Compute residuals, rotate them, and quantize them with compressing-level quantizer -> 32 compressing ids
- Add the vector to inverted multi index database, with indexing ids as key, and list of compressed residuals as value.

## Indexing process

- Rotate features, and quantize them with indexing-level quantizer -> 2 indexing ids
- Compute residuals, rotate them, and quantize them with compressing-level quantizer -> 32 compressing ids
- Add the vector to inverted multi index database, with indexing ids as key, and list of compressed residuals as value.
- This takes around 10ms per vector

## Querying process

- Rotate features, find nearest centroids with multi-sequence algorithm

## Querying process

- Rotate features, find nearest centroids with multi-sequence algorithm
- Get compressed residuals from database, to build a list of candidates

## Querying process

- Rotate features, find nearest centroids with multi-sequence algorithm
- Get compressed residuals from database, to build a list of candidates
- Rotate features again, build a look-up table for fast distance computation

## Querying process

- Rotate features, find nearest centroids with multi-sequence algorithm
- Get compressed residuals from database, to build a list of candidates
- Rotate features again, build a look-up table for fast distance computation
- Rerank the list of candidates by increasing distance to query vector



## Querying process

- Rotate features, find nearest centroids with multi-sequence algorithm
- Get compressed residuals from database, to build a list of candidates
- Rotate features again, build a look-up table for fast distance computation
- Rerank the list of candidates by increasing distance to query vector
- This takes around 500ms-1s per query\*

## Libraries used

- RocksDB (Facebook): embeddable persistent key-value store for fast storage. It stores the inverted multi index data.
- FlatBuffers (Google): serialization lib that provides access to data without parsing. Used as values in RocksDB.
- Eigen for matrix computation
- gRPC & Protocol Buffers as always

# Machines and Servers

- 1 similia-indexer machine (2 vCPUs, 1.8 GB RAM) with 3 `similia_indexer_processor`
- 2 similia-server machines (4 vCPUs, 15 GB RAM, local SSD)
- Each similia-server machine has an `inverted_multi_index_server` and a `similia_server`
- Librarian `GetSimilarImages` RPC wraps the call to `similia_server` with features retrieval/computation/storage

# Conclusion

- One can find similar images on instagram by providing to our API an url of a .jpg
- Currently 120M+ images indexed
- Fast ( $<1s$ ) \*
- Precise \*
- \*Mileage may vary



Thanks (everyone, with special thanks to Christian, Jean, and Mathieu)